


2

3D Viewing pipeline

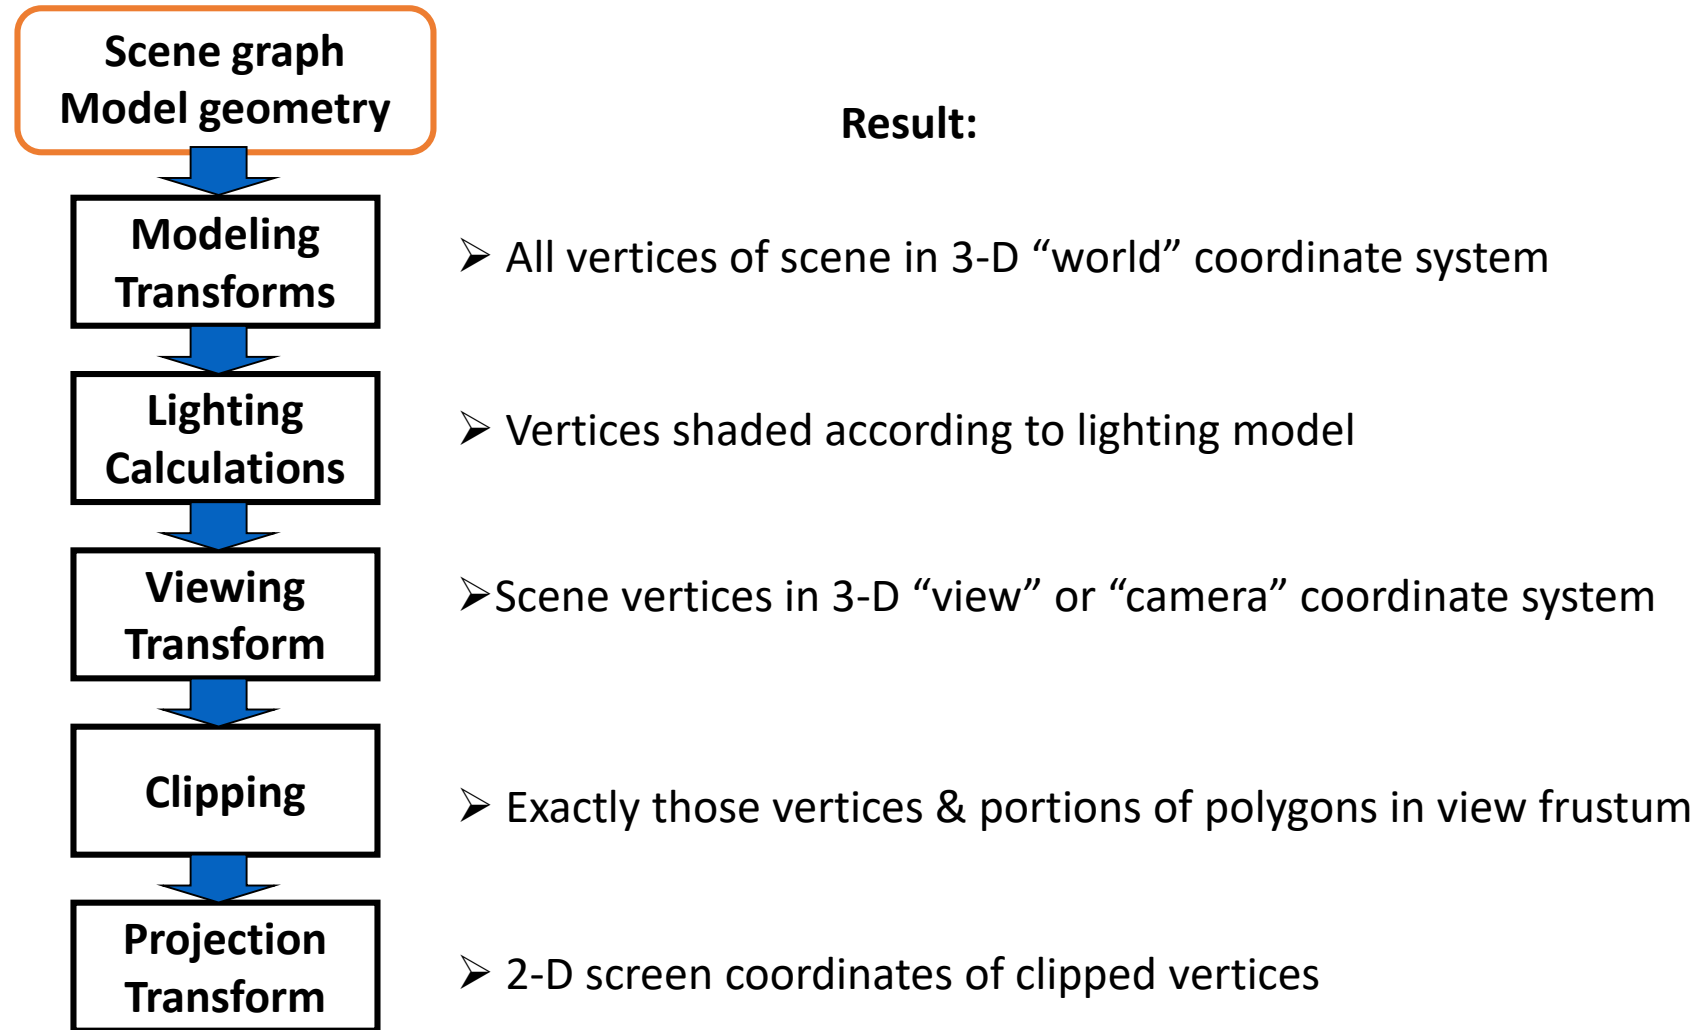
3D rendering pipeline



3D viewing pipeline

- Goal: map 3D locations, represented by (x,y,z) coordinates, to coordinates in the screen, expressed in units of pixel.
- Depends on various parameters: camera position, camera orientation, type of projection, field of view, and resolution of the screen.
- Most graphics systems do this with the following transformations:
 - Camera transformation (or eye transformation): places the camera at the origin
 - Projection transformation: projects points from camera space to normalized device coordinates (i.e. points fall in the range -1 to 1)
 - Viewport transformation (or windowing transformation): maps the image rectangle (in the range -1 to 1) to the desired rectangle in pixel coordinates

3D viewing pipeline



Modeling transformation

- Bring the object from its local coordinate system to the world coordinate system
- Scale, move, rotate objects or their parts w.r.t. each other
- Simple modeling transformation in OpenGL:
 - **`glTranslatef(GLfloat x, GLfloat y, GLfloat z);`**
Moves an object by the given *x*, *y* and *z* values.
 - **`glRotatef(GLfloat angle, GLfloat x, GLfloat y, GLfloat z);`**
Rotates an object counterclockwise about the vector from the origin through the point (*x*, *y*, *z*). The *angle* is in degrees.
 - **`glScalef(GLfloat x, GLfloat y, GLfloat z);`**
Scale each *x*, *y* and *z* coordinate of every point in the object by the corresponding argument *x*, *y* or *z*.

Transformation order

Transformation order matters.

The transformations are right-concatenated. Meaning that the following code:

```
glTranslatef(1.f,0.f,0.f);  
glRotatef(45.f,0.f,1.f,0.f);  
drawModel();
```

First rotate by 45 degrees about the Y (0,1,0) vector, and then translate by 1 unit along X.

Saving/restoring transformations

Transformations are concatenated to the right (first to be applied is the last called).

Sometimes we may want to save the current transformation state and restore it later.

`glPushMatrix()`: to save the current transformation state.


`glPopMatrix()`: to restore the last save transformation state.

Saving/restoring transformations

```
glPushMatrix(); // save the current state of transformation A
glTranslatef(1.f,0.f,0.f); // apply a couple of transformations
glRotatef(45.f,0.f,1.f,0.f);
glTranslatef(0.f,1.f,0.f);
drawModel();
glPopMatrix(); // restore to the transformation state we were in A
```


OpenGL camera

- In OpenGL, initially the object and camera frames are the same
 - Default model-view transformation is an identity
- The camera is located at origin and points in the negative z direction
- OpenGL also specifies a default view volume that is a cube with sides of length 2 centered at the origin
 - Default projection transformation is an identity



Moving the camera frame


- Once objects are positioned in world space, we need to position the scene w.r.t. the camera
- If we want to visualize object with both positive and negative z values we can either
 - Move the camera in the positive z direction
 - Translate the camera frame
 - Move the objects in the negative z direction
 - Translate the world frame
- Both of these views are equivalent and are determined by the model-view transformation
 - Want a translation (`glTranslatef (0.0f, 0.0f, -d) ;`)
 - $d > 0$

Moving the camera frame from the origin


frames after translation by $-d$

$d > 0$

default frames




(a)



(b)

Moving the camera

- We can move the camera to any desired position by a sequence of rotations and translations
- Example: side view
 - Move the camera away from origin
 - Then rotate it
 - Or: symmetrically, we can rotate the scene then move it away from the camera




LookAt function

- The GLU library contains the function `gluLookAt` to form the required model-view transformation through a simple interface
- Note the need for setting an up direction
- Can be obtained with modeling transformations
- Example: isometric view of cube aligned with axes

LookAt

```
gluLookAt(eyex, eyey, eyez,  
atx, aty, atz, upx, upy, upz);
```




Projections

- Standard projections project onto a plane
- Projectors are lines that either
 - converge at a center of projection
 - are parallel
- Such projections preserve lines
 - but not necessarily angles
- Nonplanar projections are needed for applications such as map construction


Perspective vs parallel

- Computer graphics treats all projections the same and implements them with a single pipeline
- Classical viewing developed different techniques for drawing each type of projection
- Fundamental distinction is between parallel and perspective viewing even though mathematically parallel viewing is the limit of perspective viewing

Perspective projection




Parallel projection





Orthographic projection

- Projection is parallel to one of the principal faces
- Projectors are orthogonal to the projection surface



Multi-view orthographic projection

- Typical of CAD systems
- One oblique (or sometimes perspective) view
- And three orthographic views corresponding to front, side and top




Orthographic projection

- Preserves both distances and angles
 - Shapes preserved
 - Can be used for measurements
 - Building plans
 - Manuals
- Cannot see what object really looks like because many surfaces hidden from view
 - Often we add the isometric

OpenGL orthographic projection


`glOrtho(left, right, bottom, top, near, far)`



`near` and `far` measured from camera

Perspective projection

- Projectors converge at center of projection




Perspective projection

- Objects further from viewer are projected smaller than the same sized objects closer to the viewer (*diminution*)
 - Looks realistic
- Equal distances along a line are not projected into equal distances (*non-uniform foreshortening*)
- Angles preserved only in planes parallel to the projection plane
- More difficult to construct by hand than parallel projections (but not more difficult by computer)


OpenGL perspective projection

`glFrustum(left, right, bottom, top, near, far)`



Using Field of View

- With **Frustum** it is often difficult to get the desired view
- **gluPerspective(fovy, aspect, near, far)** often provides a better interface




front plane

$$\text{aspect} = w/h$$


Orthographic projection vs perspective projection

- Consider the results of applying an orthographic projection or a perspective projection to identical cubes but at different distance along the z-axis from the camera

Orthographic projection





Perspective projection




Default projection

- In the default orthographic view, points are projected forward along the z axis onto the plane $z=0$




Viewports

- Convert from normalized device coordinates to screen coordinates (in pixel units):
`glViewport(x, y, w, h)`
- Do not have use the entire window for the image:
- Values in pixels (window coordinates)



Example: 3D cube

```
#include <GL/glut.h>
#include <stdlib.h>
GLint faces[6][4] = {{0, 1, 2, 3}, {3, 2, 6, 7},
    {7, 6, 5, 4}, {4, 5, 1, 0},
    {5, 6, 2, 1}, {7, 4, 0, 3} };
GLfloat v[8][3];
```



Vertex specification

```
void init(void)
{
    glClearColor (0.0, 0.0, 0.0, 0.0);
    v[0][0] = v[1][0] = v[2][0] = v[3][0] = -1;
    v[4][0] = v[5][0] = v[6][0] = v[7][0] = 1;
    v[0][1] = v[1][1] = v[4][1] = v[5][1] = -1;
    v[2][1] = v[3][1] = v[6][1] = v[7][1] = 1;
    v[0][2] = v[3][2] = v[4][2] = v[7][2] = -4;
    v[1][2] = v[2][2] = v[5][2] = v[6][2] = -6;
}

void keyboard(unsigned char key, int x, int y) {
    if(key=='q' || key=='Q') exit(0);
}
```

Cube in wireframe mode

```
void drawBox(void)
{
    int i;
    for (i = 0; i < 6; i++) {
        glBegin(GL_LINE_LOOP);
        glVertex3fv(&v[faces[i][0]][0]);
        glVertex3fv(&v[faces[i][1]][0]);
        glVertex3fv(&v[faces[i][2]][0]);
        glVertex3fv(&v[faces[i][3]][0]);
        glEnd();
    }
}
```



Display / reshape

```
void display(void) {
    glClear (GL_COLOR_BUFFER_BIT);
    glColor3f (1.0, 1.0, 1.0);
    drawBox();
    glFlush();
}

void reshape (int w, int h){
    glViewport (0, 0, (GLsizei) w, (GLsizei) h);
    glMatrixMode (GL_PROJECTION);
    glLoadIdentity ();
    gluPerspective(60.0, (GLfloat)w/ (GLfloat)h, 1.0, 20.0);
    glMatrixMode(GL_MODELVIEW);
    glLoadIdentity();
}
```


Main

```
int main(int argc, char** argv)
{
    glutInit(&argc, argv);
    glutInitDisplayMode (GLUT_SINGLE | GLUT_RGB);
    glutInitWindowSize (200, 200);
    glutInitWindowPosition (100, 100);
    glutCreateWindow (argv[0]);
    init ();
    glutDisplayFunc(display);
    glutReshapeFunc(reshape);
    glutKeyboardFunc(keyboard);
    glutMainLoop();
    return 0;
}
```



Using gluLookAt

```
gluLookAt (eyex, eyey, eyez, atx, aty, atz, upx, upy, upz) ;
```




```
gluLookAt (3., 3., 0., 0., 0., -5., 0., 1., 0.) ;
```

Using gluLookAt


```
void display(void) {
    glClear (GL_COLOR_BUFFER_BIT);
    glColor3f (1.0, 1.0, 1.0);
    glLoadIdentity();
    gluLookAt(3., 3., 0., 0., 0., -5., 0., 1., 0.);
    drawBox();
    glFlush ();
}

void reshape (int w, int h) {
    glViewport (0, 0, (GLsizei) w, (GLsizei) h);
    glMatrixMode(GL_PROJECTION);
    glLoadIdentity();
    gluPerspective(60.0, (GLfloat)w/h, 1.0, 40.0);
    glMatrixMode (GL_MODELVIEW);
}
```

Results



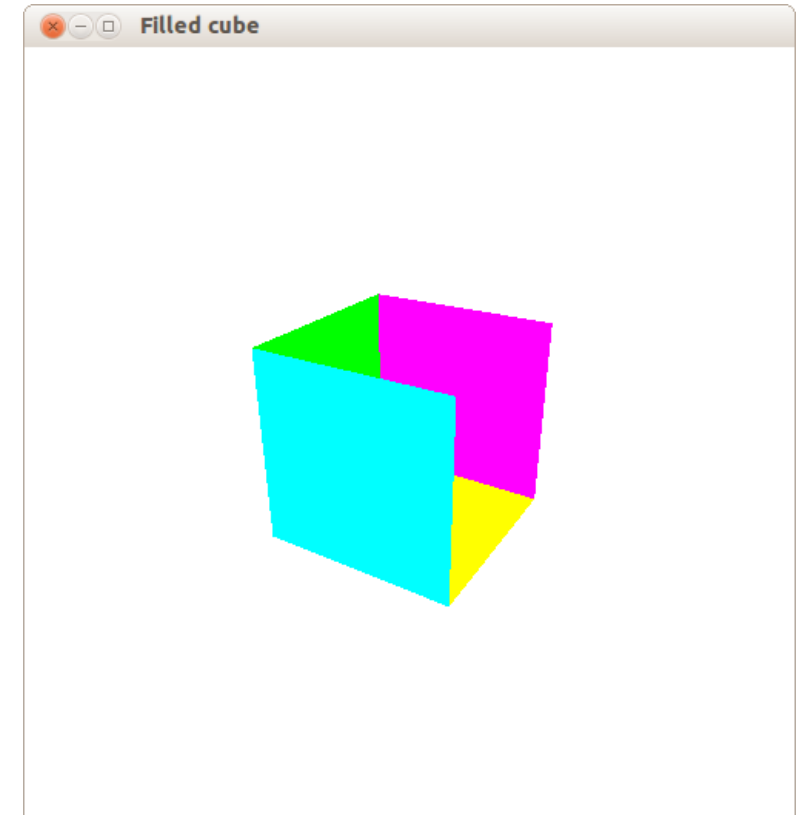
```
gluLookAt(3.0, 3.0, 0.0,  
          0.0, 0.0, -5.0,  
          0.0, 1.0, 0.0);
```



```
gluLookAt(2.0, 4.0, -3.0,  
          0.0, 0.0, -5.0,  
          0.0, 1.0, 0.0)
```

Rendering filled cube (hidden surface removal)

- Let us try to fill the cube faces by replacing: `glBegin(GL_LINE_LOOP)` by: `glBegin(GL_POLYGON)`
- and by using different colors for each face.
- The result is not satisfactory




Rendering filled cube


```
// closest triangle  
glColor3f(1.0f, 0.0f, 0.0f);  
glBegin(GL_TRIANGLES);  
    glVertex3f(0.25f, 0.25f, 0.0f);  
    glVertex3f(0.75f, 0.25f, 0.0f);  
    glVertex3f(0.75f, 0.75f, 0.0f);  
glEnd();
```

```
// farthest triangle  
glColor3f(0.0f, 1.0f, 0.0f);  
glBegin(GL_TRIANGLES);  
    glVertex3f(0.25f, 0.25f, -0.5f);  
    glVertex3f(0.75f, 0.25f, -0.5f);  
    glVertex3f(0.25f, 0.75f, -0.5f);  
glEnd();
```

Expected



Obtained



Explanation

Polygons are sent to the pipeline and rendered *in order*.

In the example:

- First defined triangle is the red one
- Then the green one
- Therefore the obtained result (independently of the position of the triangles w.r.t the camera)

Rule: the last triangle to be rendered is the one that will be seen

Hidden surface removal

- This problem is called: *Hidden Surface Elimination*
- We want to see the triangles closest to the camera (and not the last defined triangles)
- Solution: use an additional buffer, called the *depth-buffer* (also sometimes z-buffer) that stores depth value for each pixel and perform tests (*depth-test*) to determine what to draw

Depth-buffer and depth-test

- *Depth-buffer* stores depth values for each pixel.
- The idea is then to stop writing a fragment if it is in the back of the current pixel. This is called the *depth-test* (sometimes also called z-buffer algorithm).

Depth-buffer and depth-test in OpenGL

To activate depth-testing:

```
glEnable(GL_DEPTH_TEST)
```

The corresponding `glDisable` causes depth-testing to stop.

Before doing depth-tests, we need to initialize the depth-buffer

In order to clear the depth buffer, we use `glClear` (like for the color buffer):

```
glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);
```

Steps to use depth-test in OpenGL

1. Create a GL rendering context with depth-buffer enabled:

```
glutInitDisplayMode(GLUT_SINGLE | GLUT_RGBA |  
GLUT_DEPTH) ;
```

2. Clear the depth-buffer before usage:

```
glClear(GL_COLOR_BUFFER_BIT |  
GL_DEPTH_BUFFER_BIT) ;
```

3. Enable the depth-test: `glEnable(GL_DEPTH_TEST) ;`

Example

```
static GLshort g_use_depth_buffer = 0;

void initGL() {
    glClearColor(1.0, 1.0, 1.0, 1.0);

    glMatrixMode(GL_PROJECTION);
    glLoadIdentity();
    glOrtho(0.0, 1.0, 0.0, 1.0, -1.0, 1.0);

    glMatrixMode(GL_MODELVIEW);
    glLoadIdentity();
}
```

Example (continued)

```
void handleDisplay() {
    glClear (GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);

    if (g_use_depth_buffer) {
        glEnable(GL_DEPTH_TEST);
        printf("Depth-buffer enabled\n");
    } else {
        printf("Depth-buffer disabled\n");
        glDisable(GL_DEPTH_TEST);
    }

    // closest triangle
    glColor3f(1.0f, 0.0f, 0.0f);
    glBegin(GL_TRIANGLES);
    glVertex3f(0.25f, 0.25f, 0.0f);
    glVertex3f(0.75f, 0.25f, 0.0f);
    glVertex3f(0.75f, 0.75f, 0.0f);
    glEnd();

    // ...
}
```

```
void handleDisplay() {
    // ...

    // farthest triangle
    glColor3f(0.0f, 1.0f, 0.0f);
    glBegin(GL_TRIANGLES);
        glVertex3f(0.25f, 0.25f, -0.5f);
        glVertex3f(0.75f, 0.25f, -0.5f);
        glVertex3f(0.25f, 0.75f, -0.5f);
    glEnd();

    glFlush();
}
```

Example (continued)

```
void
handleKeyboardEvents(unsigned char key, int x, int y)
{
    if (key == 'q' || key == 'Q') {
        exit(0);
    }

    if (key == 'z') {
        g_use_depth_buffer = 1 - g_use_depth_buffer;
        glutPostRedisplay();
        return;
    }
}
```

```
int main(int argc, char** argv) {
    glutInit(&argc, argv);
    glutInitDisplayMode(GLUT_SINGLE | GLUT_RGBA | GLUT_DEPTH);
    glutInitWindowSize(500, 500);
    glutInitWindowPosition(100, 100);
    glutCreateWindow("Depth-buffer");

    initGL();

    glutDisplayFunc(handleDisplay);
    glutKeyboardFunc(handleKeyboardEvents);

    glutMainLoop();

    return EXIT_SUCCESS;
}
```

Result

- Let us return to the filled cube example, this time using the depth-buffer and a depth-test:

